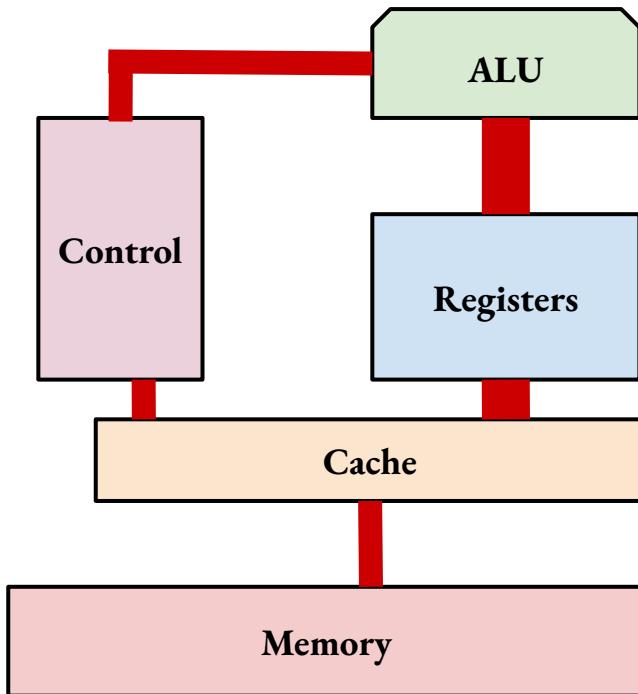


Scalable Pipeline Parallelism on FPGAs with High-Level Synthesis

May 28, 2025

Load Store Architecture -

Load - Store (“von Neumann”)



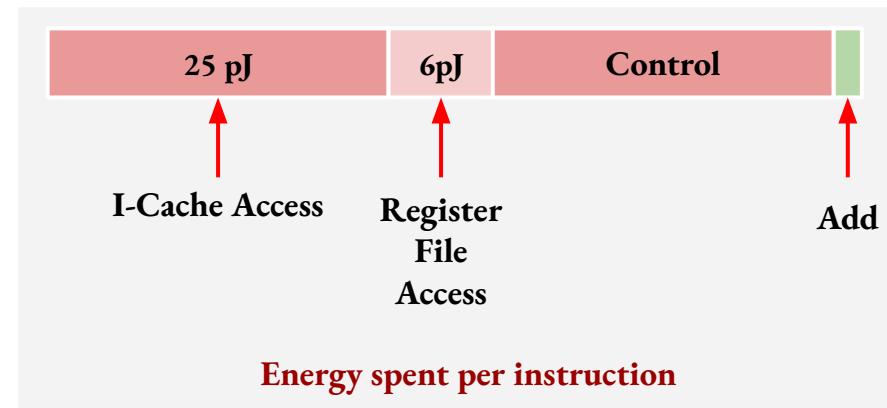
1: ld a, r1
2: ld b, r2
3: add r1, r2, r1
4: st c, r1

Addition of two numbers

32-bit FP Add: 0.9 pJ
32-bit FP Mul: 3.2 pJ

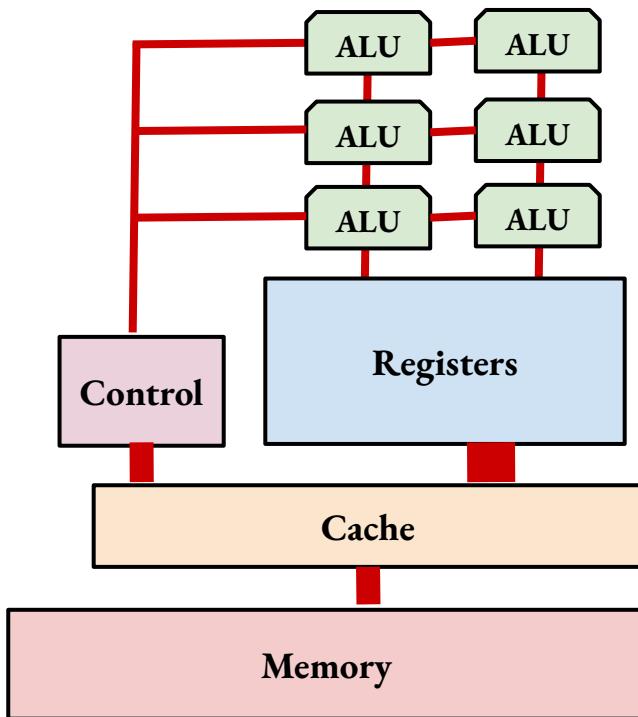
2x32 bit from L1 (8 kiB): 10 pJ
2x32 bit from L2 (1MiB): 100 pJ
2x32 bit from DRAM: 1.3 nJ

*Electrical data movement loss ~
resistance * length / cr.area*

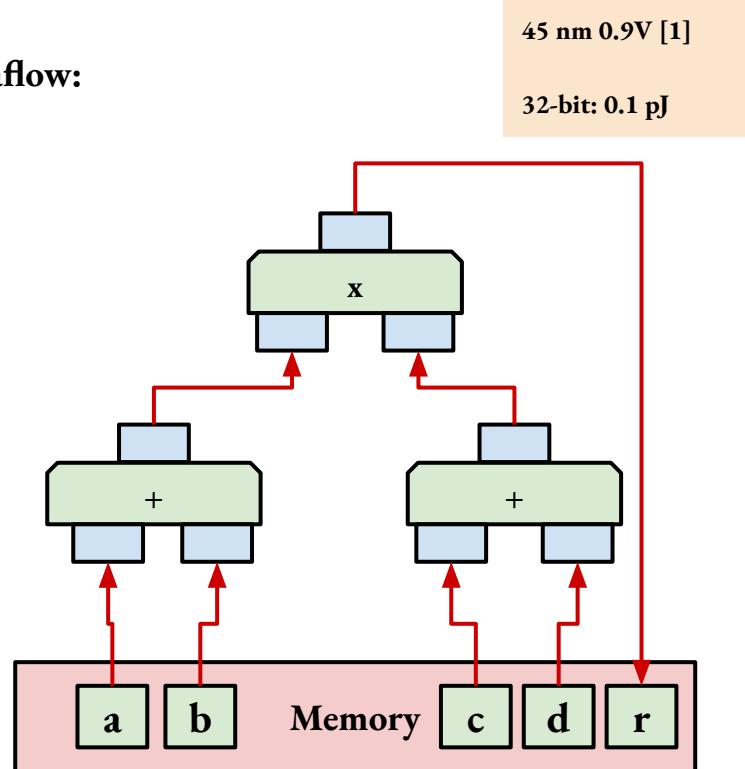


Load Store Architecture -

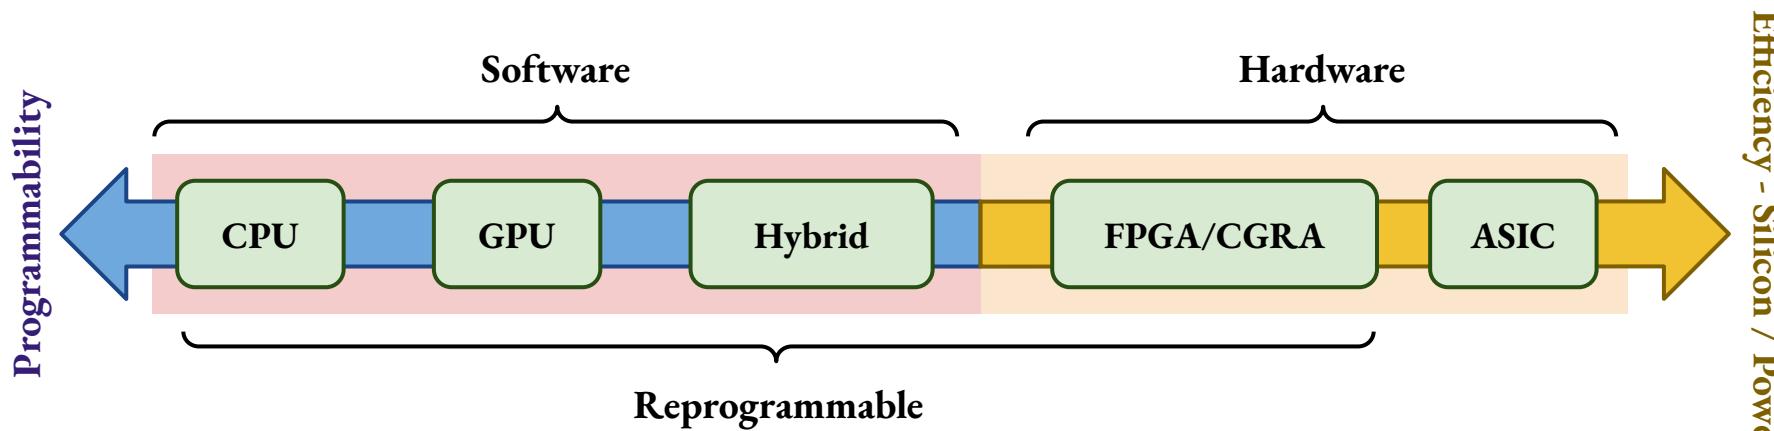
SIMD / SIMD



Dataflow:

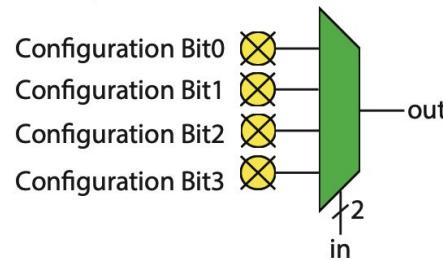
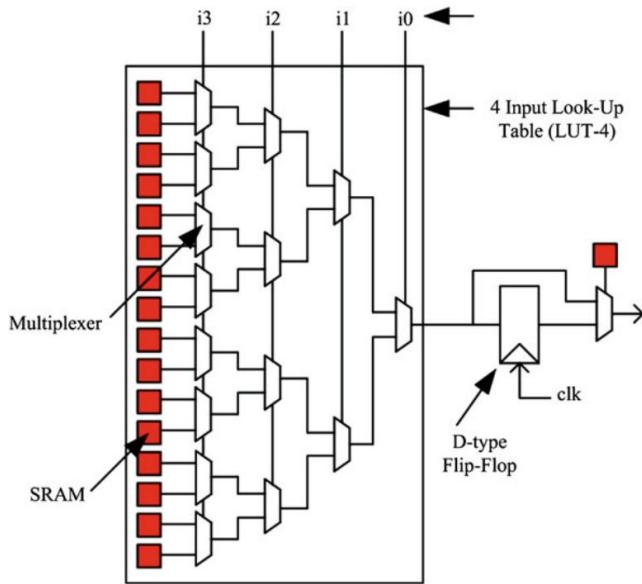


Accelerators for High Performance Computing -



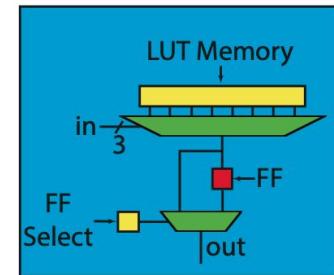
Reducing data movement is a challenge for obtaining higher performance

Field Programmable Gate Arrays (FPGA)



in[1]	in[0]	out
0	0	0
0	1	0
1	0	0
1	1	1

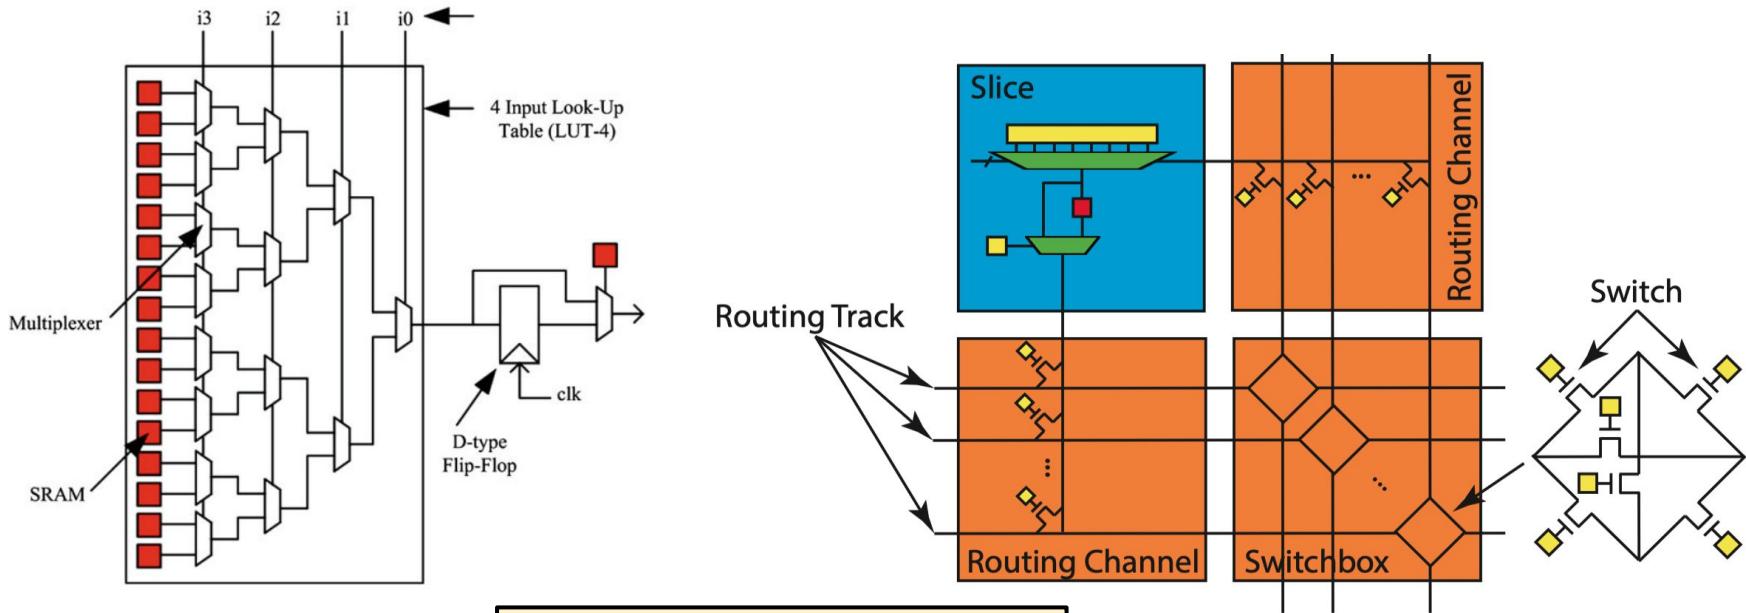
out = in[1] & in[0]



Configuration logic block

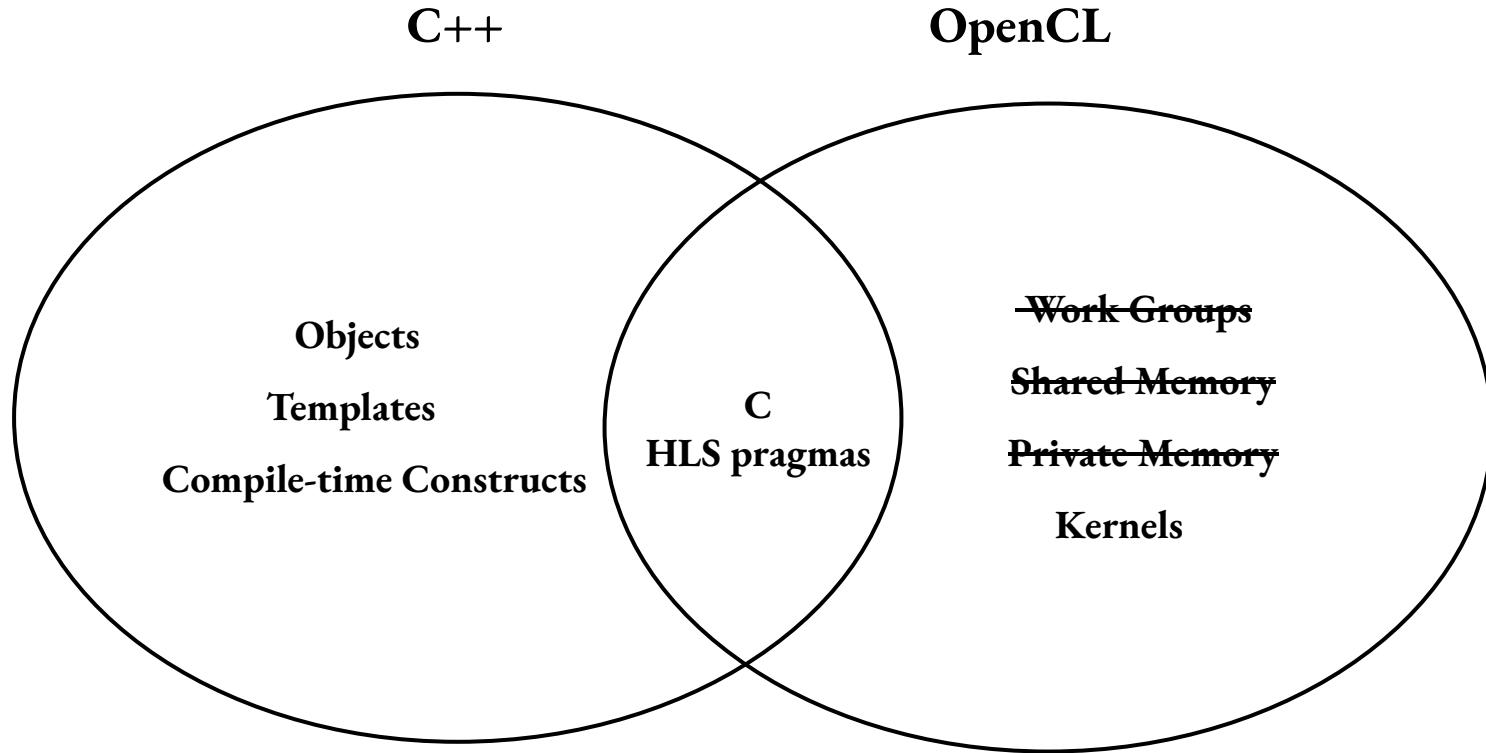
Implementing a circuit on the FPGA is about filling the **red** memory cells / **configuration** bits.

Field Programmable Gate Arrays (FPGA)



Implementing a circuit on the FPGA is about filling the **red** memory cells / **configuration** bits.

Programming FPGAs with HLS

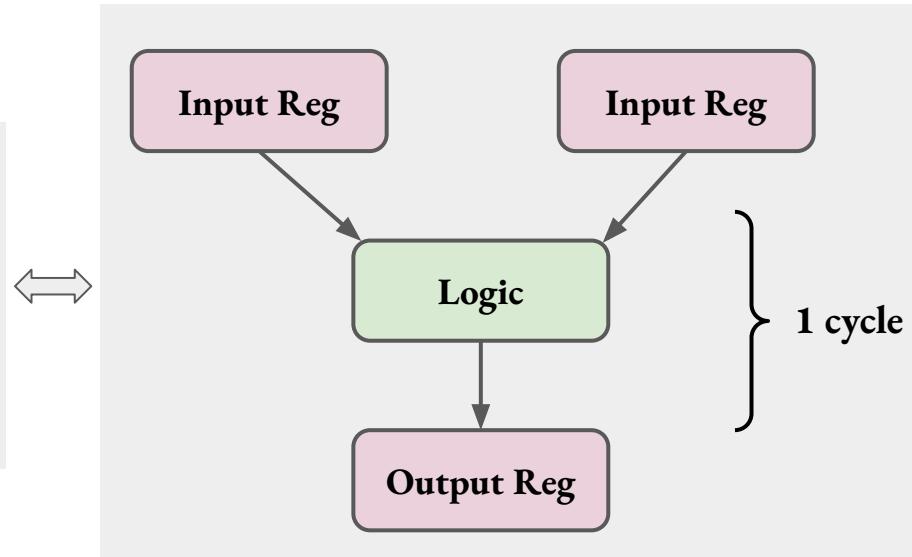


End Goal: Perform compute in every piece of available logic in every cycle.

Register Transfer Level

```
int c = a + b;
```

```
always @(posedge clk)
  if (start) begin
    c <= a + b;
  end
```



Register Transfer Level

float c = a + b;

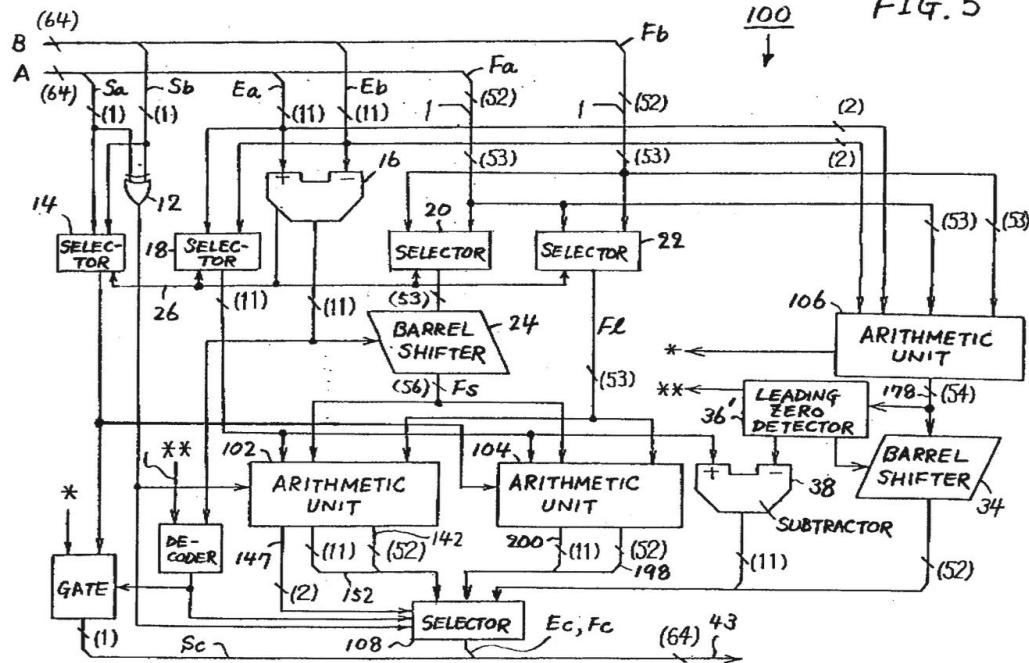
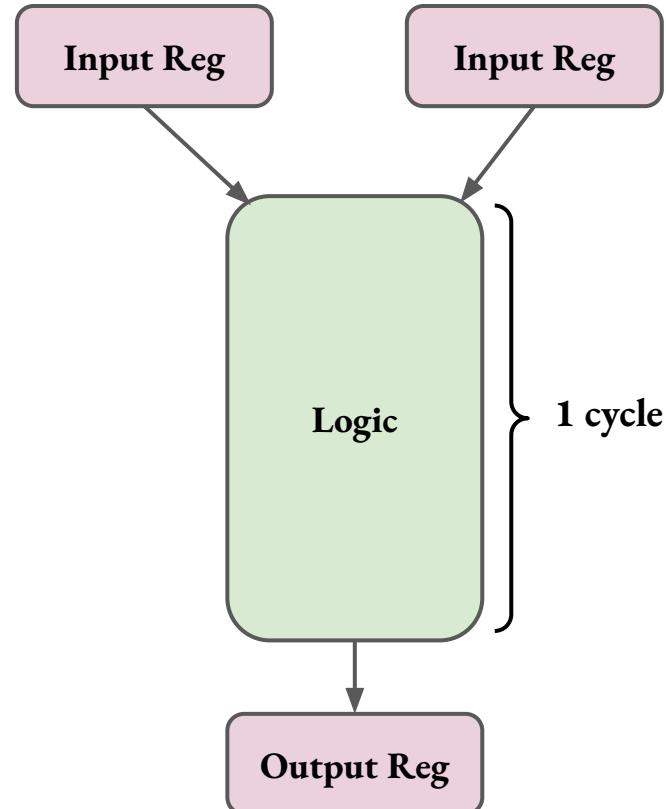


FIG. 5



Register Transfer Level

```
float c = a + b;
```

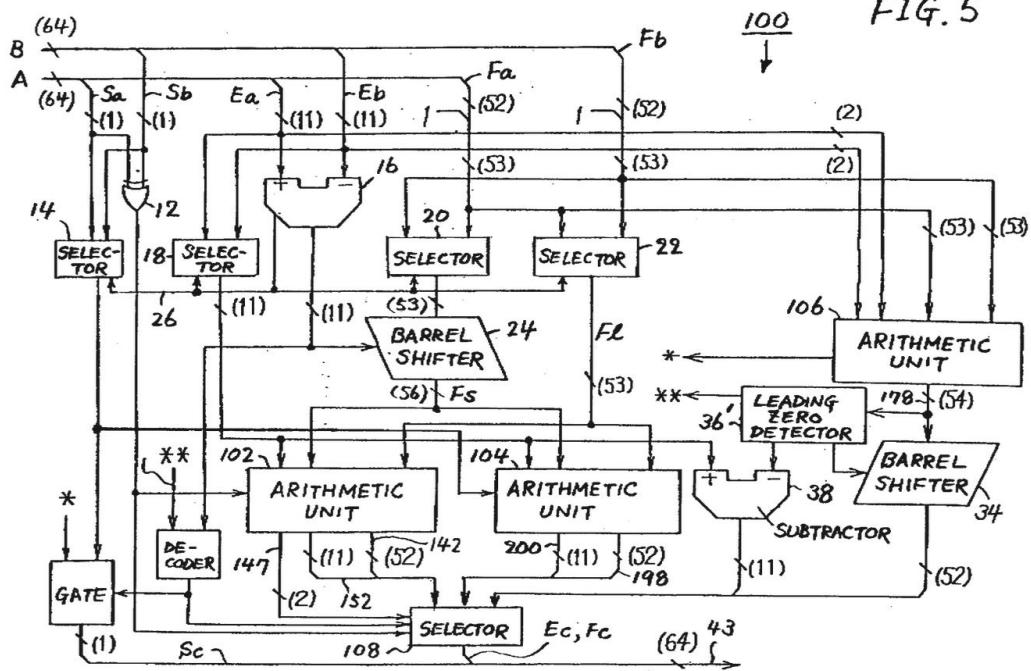
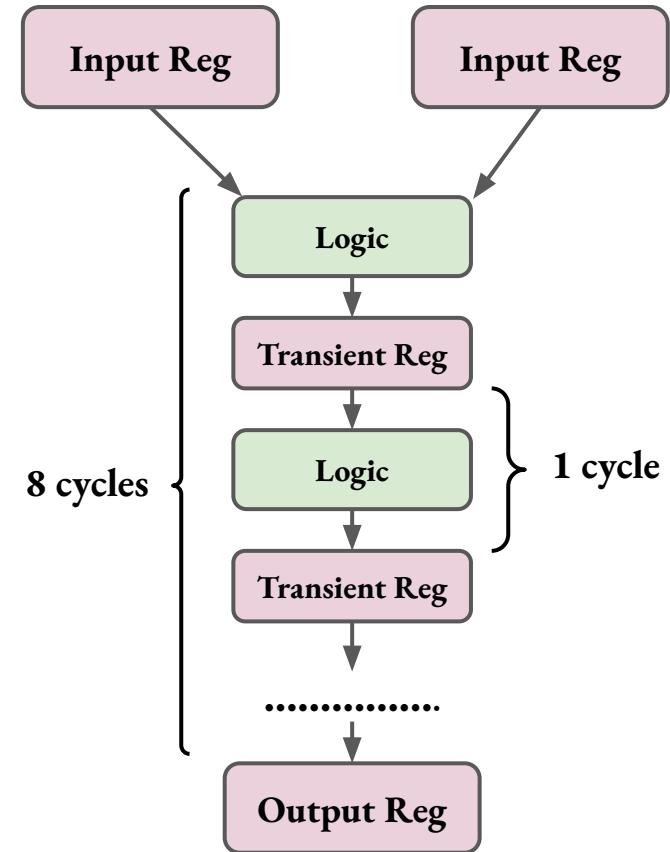


FIG. 5



Register Transfer Level

float c = a + b;

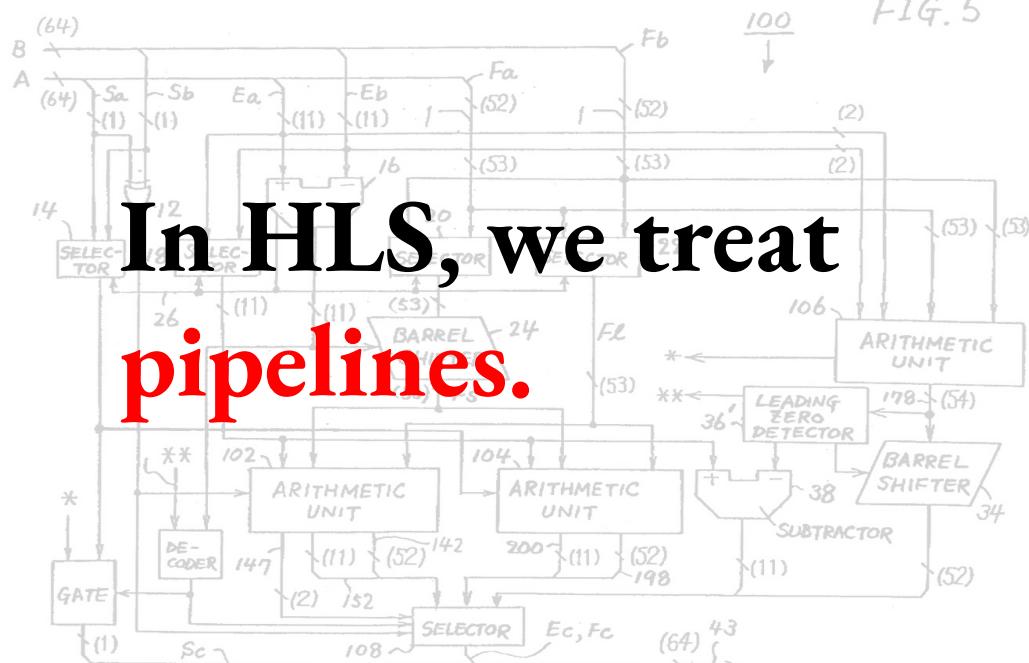
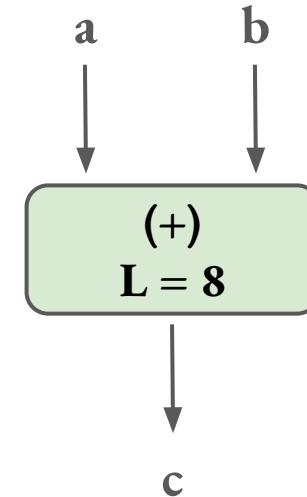
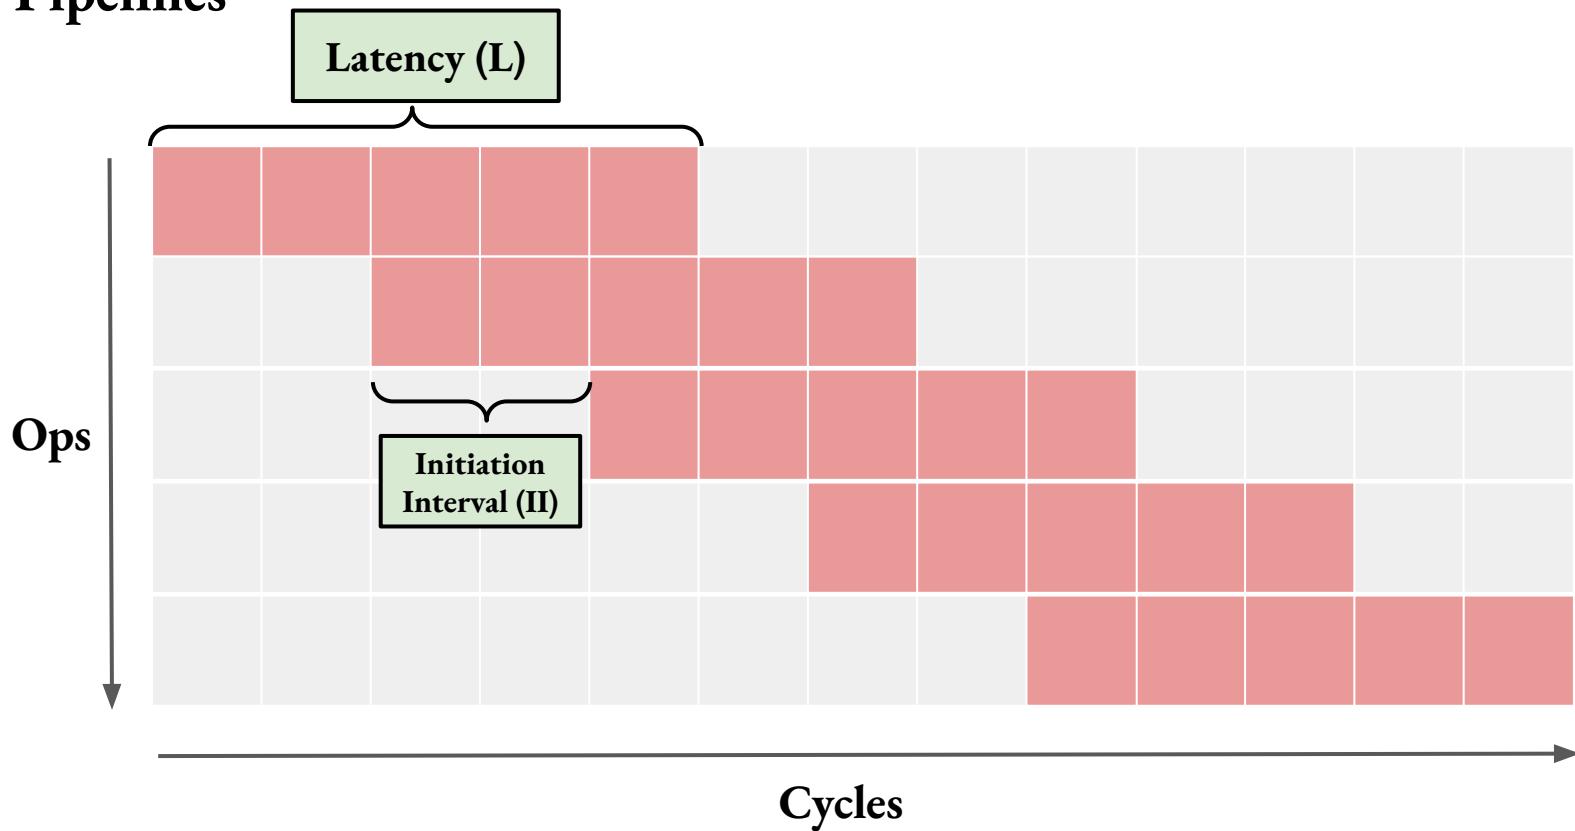


FIG. 5

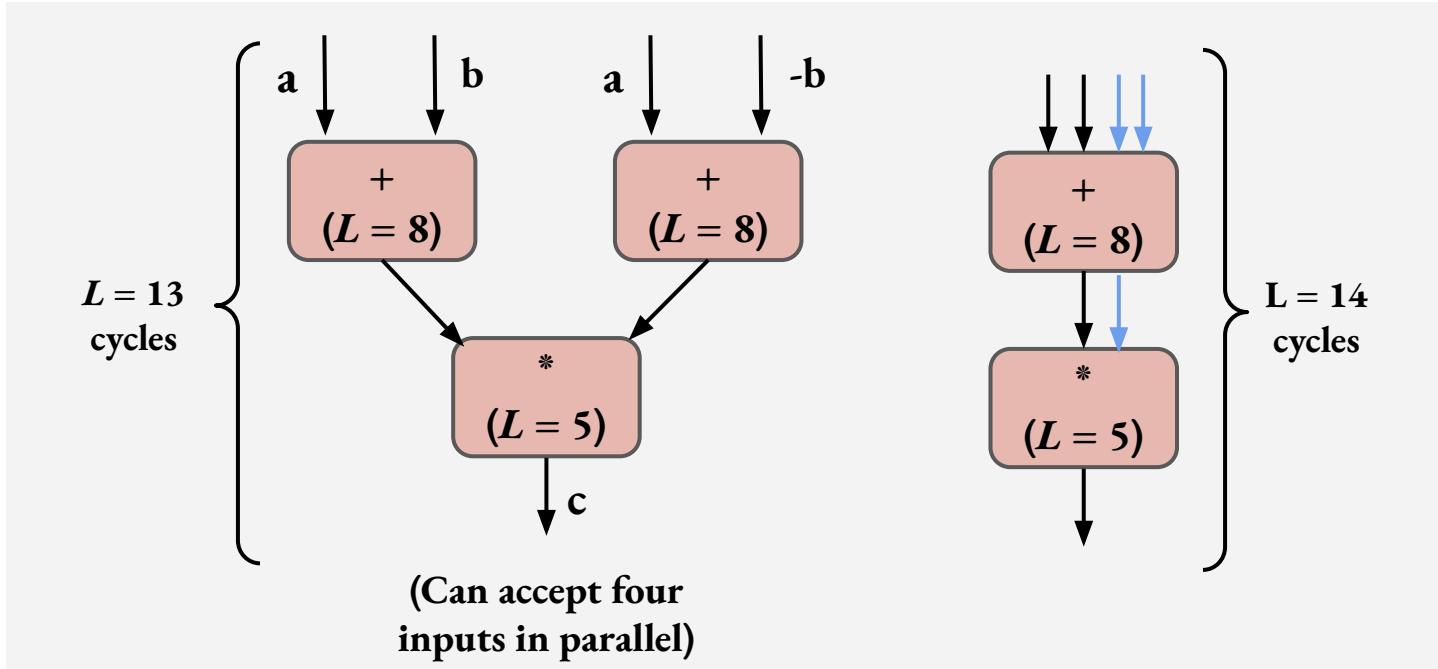


Implies $L - 1$ internal pipeline stages.

Pipelines



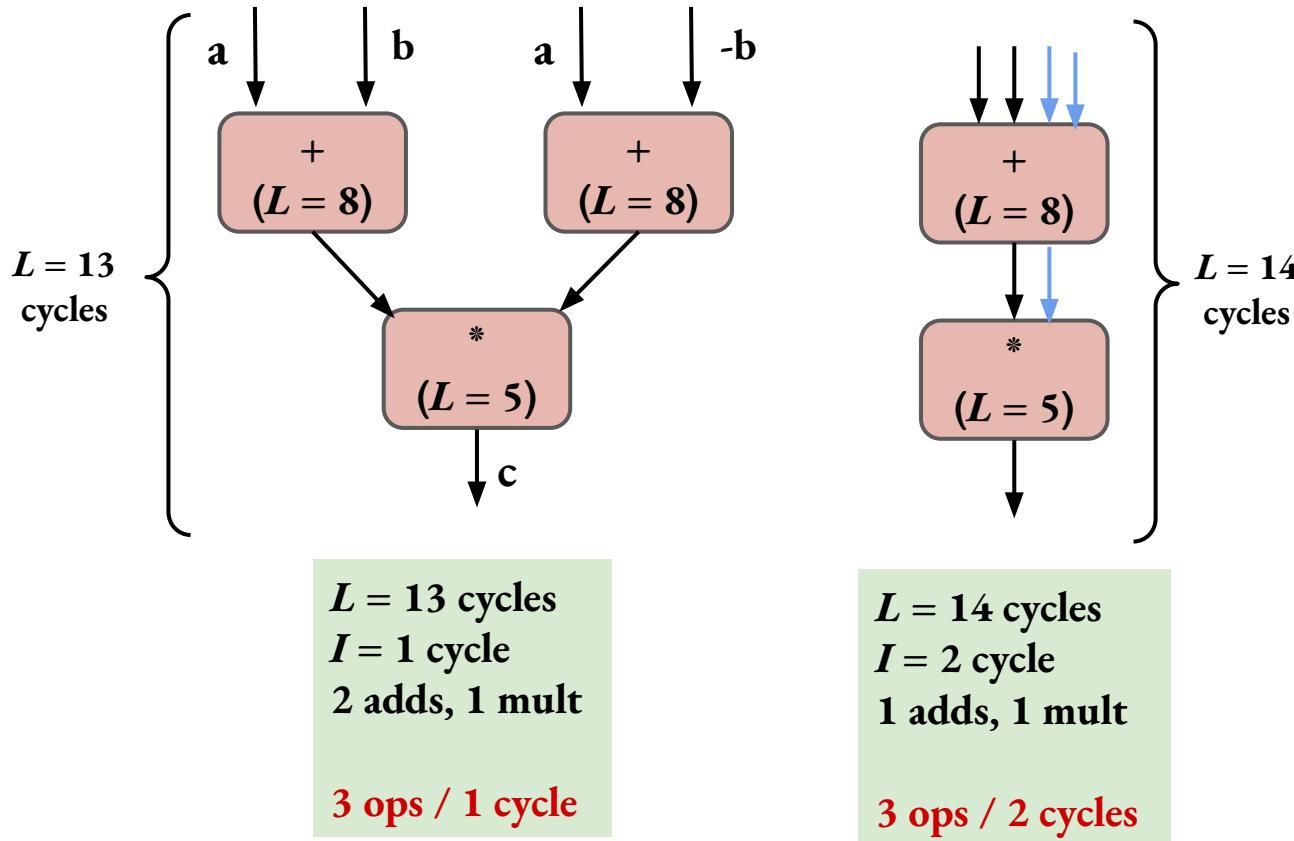
Multiple floating-point operations



```
float c = (a + b) * (a - b);
```

Two or more ways to implement this

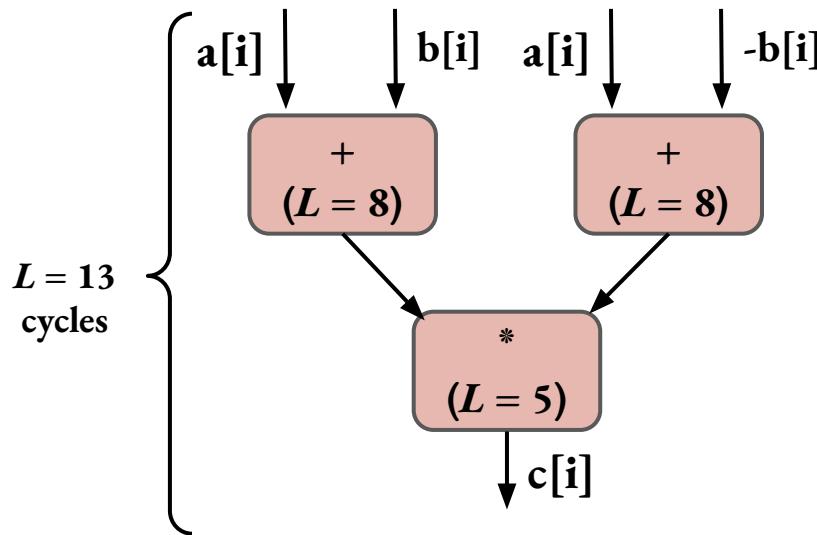
Multiple floating-point operations



In addition to **latency** (L), we have **initiation interval** (I):

1. No of cycles before we can accept new inputs.
2. Inverse throughput of the pipeline.
3. Factor of slowdown of the application.

Multiple floating-point operations



$L = 13$ cycles
 $I = 1$ cycle
2 adds, 1 mult
3 ops / 1 cycle

```
for (int i = 0; i < N; ++i) {  
    #pragma HLS PIPELINE II=1  
    c[i] = (a[i] + b[i]) *  
           (a[i] - b[i]);  
}
```

1 iteration	$13 + 1 = 14$ cycles
10 iteration	$13 + 10 = 23$ cycles

N iteration	$13 + N$ cycles
---------------------------------	-----------------

(Loops iterations affect the running time additively)